

Final Exam, CPSC 8420, Spring 2022

Huang, Gangtong

Due 05/06/2022, Friday, 11:59PM EST

Problem 1 [15 pts]

Consider the elastic-net optimization problem:

$$\min_{\beta} \|\mathbf{y} - \mathbf{X}\beta\|^2 + \lambda[\alpha\|\beta\|_2^2 + (1 - \alpha)\|\beta\|_1]. \quad (1)$$

1. Show the objective can be reformulated into a lasso problem, with a slightly different \mathbf{X}, \mathbf{y} .

Define: $\mathbf{X}_1 = (1 + \lambda\alpha) \begin{pmatrix} \mathbf{X} \\ \sqrt{\lambda\alpha}\mathbf{I}_p \end{pmatrix}$, $\mathbf{y}_1 = \begin{pmatrix} \mathbf{y} \\ 0 \end{pmatrix}$, and $\beta_1 = \sqrt{1 + \lambda\alpha}\beta$ then:

$$\begin{aligned} \mathbf{X}_1^T \mathbf{X}_1 &= (1 + \lambda\alpha)(\mathbf{X}^T \mathbf{X} + \lambda\alpha \mathbf{I}) \\ &= (1 + \lambda\alpha)(\mathbf{X}^T \mathbf{X} + \lambda\alpha) \end{aligned} \quad (2)$$

, and

$$\begin{aligned} \mathbf{y}_1^T \mathbf{y}_1 &= \mathbf{y}^T \mathbf{y} + 0 \\ &= \mathbf{y}^T \mathbf{y} \end{aligned} \quad (3)$$

Define $\mathbf{L}(\mathbf{y}_1, \mathbf{X}_1, \beta_1) = \|\mathbf{y}_1 - \mathbf{X}_1\beta_1\|^2 + \frac{\lambda(1-\alpha)}{1+\lambda\alpha}\|\beta_1\|_1$. Formulata a Lasso problem using $\mathbf{X}_1, \mathbf{y}_1, \beta_1$:

$$\min_{\beta_1} \mathbf{L}(\mathbf{y}_1, \mathbf{X}_1, \beta_1) \quad (4)$$

, where:

$$\begin{aligned} \mathbf{L}(\mathbf{y}_1, \mathbf{X}_1, \beta_1) &= \|\mathbf{y}_1 - \mathbf{X}_1\beta_1\|^2 + \frac{\lambda(1-\alpha)}{1+\lambda\alpha}\|\beta_1\|_1 \\ &= (\mathbf{y}_1 - \mathbf{X}_1\beta_1)^T (\mathbf{y}_1 - \mathbf{X}_1\beta_1) + \frac{\lambda(1-\alpha)}{1+\lambda\alpha}\|\beta_1\|_1 \\ &= \mathbf{y}_1^T \mathbf{y}_1 - \beta_1^T \mathbf{X}_1^T \mathbf{y}_1 - \mathbf{y}_1^T \mathbf{X}_1 \beta_1 + \beta_1^T \mathbf{X}_1^T \mathbf{X}_1 \beta_1 + \frac{\lambda(1-\alpha)}{1+\lambda\alpha}\|\beta_1\|_1 \end{aligned} \quad (5)$$

Plug eqns. (2), (3) into eqn. (4), we will arrive at:

$$\mathbf{L}(\mathbf{y}_1, \mathbf{X}_1, \beta_1) = \|\mathbf{y}_1 - \mathbf{X}_1\beta_1\|^2 + \frac{\lambda(1-\alpha)}{1+\lambda\alpha}\|\beta_1\|_1 \quad (6)$$

, and the optimization problem in eqn. (4) becomes:

$$\min_{\beta_1} \|\mathbf{y}_1 - \mathbf{X}_1\beta_1\|^2 + \frac{\lambda(1-\alpha)}{1+\lambda\alpha} \|\beta_1\|_1 \quad (7)$$

, where the only difference with the original optimization problem is the variable β_1 . Since $\beta_1 = \sqrt{1+\lambda\alpha}\beta$, the optimization problem in eqn. (7) is equivalent to:

$$\min_{\beta} \|\mathbf{y}_1 - \mathbf{X}_1\beta_1\|^2 + \frac{\lambda(1-\alpha)}{1+\lambda\alpha} \|\beta_1\|_1 \quad (8)$$

, which is the original objective function of the elastic net problem.

2. If we fix $\alpha = .5$, please derive the closed solution by making use of alternating minimization that each time we fix the rest by optimizing one single element in β . You need randomly generate \mathbf{X}, \mathbf{y} and initialize β_0 , and show the objective decreases monotonically with updates.

For the Lasso problem $\min_{\beta_1} \|\mathbf{y}_1 - \mathbf{X}_1\beta_1\|^2 + \frac{\lambda(1-\alpha)}{1+\lambda\alpha} \|\beta_1\|_1$, the closed solution of the i -th element of β_1 is:

$$\beta_{1i} = \text{signum}(\beta_{1i}^{LS}) (\|\beta_{1i}^{LS}\|_1 - \frac{\lambda(1-\alpha)}{1+\lambda\alpha})^+ \quad (9)$$

, where $\beta_{1i}^{LS} = X_1^T y_1$ is the solution of the vanilla least square problem.

Before the each time $\beta_{1i}^{i+1} - \beta_{1i}^i$ is evaluated against the tolerance, each component β_{1i} is optimized once. Fig. 1 shows the decrease of the objective function with the optimization of each single β_{1i} .




Figure 1: Decrease of $\mathbf{L}(\mathbf{y}_1, \mathbf{X}_1, \beta_1) = \|\mathbf{y}_1 - \mathbf{X}_1\beta_1\|^2 + \frac{\lambda(1-\alpha)}{1+\lambda\alpha}\|\beta_1\|_1$ with gradient descent w.r.t. each component of β_1 .

Codes (MatLab):

```

n = 3;
m = 4;

lambda = 10;
alpha = 1/2;
gamma = lambda*(1-alpha)/(1+lambda*alpha)

X = rand(m,n);
y = 5 * rand(m,1);

Xcent = zscore(X);

X_1 = (1 + lambda*alpha) * [Xcent; eye(n)];
y_1 = [y; zeros(n,1)];

%% Output L in each step
mat_obj = lassoAlg_step(X_1, y_1, gamma)
l_obj = reshape(mat_obj,1,[])

```

```

plot(l_obj)
xlabel('Iterations (on beta_i)')
ylabel('Objective')

%%
lassoAlg(X_1, y_1, gamma);
%% Lasso Optimization Algorithm %%
% inputs: A (nxd matrix), y (nx1 vector), lam (scalar)
% return: xh (dx1 vector)

function xh = lassoAlg(A,y,lam)
xnew = rand(size(A,2),1); % "initial guess"
xold = xnew+ones(size(xnew)); % used zeros so the while loop initiates
loss = xnew - xold;
thresh = 10e-3; % threshold value for optimization

while norm(loss) > thresh
xold = xnew; % need to store the previous iteration of xh
for i = 1:length(xnew)
a = A(:,i); % get column of A
p = (norm(a,2))^2;
% from notes: -t = sum(aj*xj) - y for all j != i
% i.e., sum(aj*xj) - ai*xi - y (my interpretation)
% hence t = (above) * -1
% want to be sure this the correct definition of t?
t = a*xnew(i) + y - A*xnew;
q = a'*t;
% update xi
xnew(i) = (1/p) * sign(q) * max(abs(q)-lam, 0);
end
loss = xnew - xold; % update loss
end
xh = xnew;
end

%% Lasso Optimization Algorithm %%
% inputs: A (nxd matrix), y (nx1 vector), lam (scalar)
% return: xh (dx1 vector)

function xh = lassoAlg_step(A,y,lam)
xnew = rand(size(A,2),1); % "initial guess"
xold = xnew+ones(size(xnew)); % used zeros so the while loop initiates
loss = xnew - xold;
thresh = 10e-3; % threshold value for optimization
xh = [];

```

```

while norm(loss) > thresh
xold = xnew;    % need to store the previous iteration of xh
tmp = [];
for i = 1:length(xnew)
a = A(:,i);    % get column of A
p = (norm(a,2))^2;
% from notes: -t = sum(aj*xj) - y for all j != i
% i.e., sum(aj*xj) - ai*xi - y (my interpretation)
% hence t = (above) * -1
% want to be sure this the correct definition of t?
t = a*xnew(i) + y - A*xnew;
q = a'*t;
% update xi
xnew(i) = (1/p) * sign(q) * max(abs(q)-lam, 0);
obj = norm(y - A * xnew, 2) + lam*(1-1/2)/(1+lam*1/2)*norm(xnew,1);
tmp = [tmp, obj];
end
loss = xnew - xold;    % update loss
%obj = norm(y - A * xnew, 2) + lam*(1-1/2)/(1+lam*1/2)*norm(xnew,1);
%xh = [xh, obj];
xh = [xh; tmp];
end
end

```

You may input your answers here. \LaTeX version submission is encouraged.

Problem 2 [15 pts]

- For PCA, the loading vectors can be directly computed from the q columns of \mathbf{U} where $[\mathbf{U}, \mathbf{S}, \mathbf{U}] = \text{svd}(\mathbf{X}^T \mathbf{X})$, please show that any $[\pm \mathbf{u}_1, \pm \mathbf{u}_2, \dots, \pm \mathbf{u}_q]$ will be equivalent to $[\mathbf{u}_1, \mathbf{u}_2, \dots, \mathbf{u}_q]$ in terms of the same variance while satisfying the orthonormality constraint.

In PCA, the variance of the i -th principal component (PC) is:

$$\begin{aligned} \text{Var}(u_i) &= \mathbf{u}_i^T \mathbf{X}^T \mathbf{X} \mathbf{u}_i \\ &= \mathbf{u}_i^T s_i \mathbf{u}_i \\ &= s_i \mathbf{u}_i^T \mathbf{u}_i \\ &= s_i \end{aligned} \tag{10}$$

, where s_i is the i -th singular value of X .

Substituting u_i with $-u_i$, we have:

$$\begin{aligned} \text{Var}(-u_i) &= (-\mathbf{u}_i^T) \mathbf{X}^T \mathbf{X} (-\mathbf{u}_i) \\ &= (-\mathbf{u}_i^T) s_i (-\mathbf{u}_i) \\ &= s_i \mathbf{u}_i^T \mathbf{u}_i \\ &= s_i \end{aligned} \tag{11}$$

, which is identical to the case of u_i . Also when $i \neq j$,

$$\begin{aligned} (-u_i)^T u_j &= -u_i^T u_j \\ &= 0 \end{aligned} \tag{12}$$

still satisfies orthonormality.

So $\text{Var}(u_i) = \text{Var}(-u_i)$, and any sign combination of $[\mathbf{u}_1, \mathbf{u}_2, \dots, \mathbf{u}_q]$ preserves the variance as well as orthonormality.

- We consider the case when original dimensionality of the data is much larger than the number of samples $d \gg m$ ($\mathbf{X} \in \mathbb{R}^{m \times d}$). What's the complexity of obtaining the optimal solution of PCA via Singular Value Decomposition? Please consider a more efficient solution by considering the relationships of eigenvalues/eigenvectors between $\mathbf{X}^T \mathbf{X}$ and $\mathbf{X} \mathbf{X}^T$.

Since $\mathbf{X} \in \mathbb{R}^{m \times d}$, the complexity of $X^T X$ is: $O(d \times d \times m) = O(d^2 m)$. Since $X^T X \in \mathbb{R}^{d \times d}$ is a square matrix, the complexity of $\text{SVD}(X^T X)$ is $O(d^3)$. Therefore, the complexity of PCA on $X^T X$ is $O(d^2 m + d^3)$.

We know that the non-zero eigenvalues of $X^T X$ and $X X^T$ are the same, and hence the non-zero singular values. Similar to $\text{PCA}(X^T X)$, we know that the time complexity of the PCA of $X X^T$ is $O(m^2 d + m^3)$, which is smaller than $O(d^2 m + d^3)$ when $d \gg m$. Therefore, to reduce computational cost, we can perform $\text{PCA}(X X^T)$ instead of $\text{PCA}(X^T X)$.

Problem 3 [10 pts]

Assume that in a community, there are 10% people suffer from COVID. Assume 80% of the patients come to breathing difficulty while 25% of those free from COVID also have symptoms of shortness of breath. Now please determine that if one has breathing difficulty, what's his/her probability to get COVID? (*hint*: you may consider Naive Bayes)

The conditional probability of getting shortness of breath B given getting Covid C is:

$$P(B|C) = 0.8 \quad (13)$$

The probability of shortness of breath without Covid (nC):

$$P(B|nC) = 0.25 \quad (14)$$

So the probability of getting shortness of breath is:

$$\begin{aligned} P(B) &= P(B|C)P(C) + P(B|nC)P(nC) \\ &= P(B|C)P(C) + P(B|nC)[1 - P(C)] \\ &= 0.8 \times 0.1 + 0.25 \times 0.9 \\ &= 0.305 \end{aligned} \quad (15)$$

The probability of getting Covid:

$$P(C) = 0.1 \quad (16)$$

From Naive Bayesian, the probability of Covid given shortness of breath is:

$$\begin{aligned} P(C|B) &= \frac{P(B|C)P(C)}{P(B)} \\ &= \frac{0.8 \times 0.1}{0.305} \\ &= 0.262 \end{aligned} \quad (17)$$

Problem 4 [20 pts]

Recall the objective for RatioCut: $RatioCut(A_1, A_2, \dots, A_k) = \frac{1}{2} \sum_{i=1}^k \frac{W(A_i, \bar{A}_i)}{|A_i|}$. If we introduce indicator

vector: $h_j \in \{h_1, h_2, \dots, h_k\}, j \in [1, k]$, for any vector $h_j \in R^n$, we define: $h_{ij} = \begin{cases} 0 & v_i \notin A_j \\ \frac{1}{\sqrt{|A_j|}} & v_i \in A_j \end{cases}$,

we can prove: $h_i^T L h_i = \frac{cut(A_i, \bar{A}_i)}{|A_i|}$, and therefore:

$$RatioCut(A_1, A_2, \dots, A_k) = \sum_{i=1}^k h_i^T L h_i = \sum_{i=1}^k (H^T L H)_{ii} = tr(H^T L H), \quad (18)$$

thus we relax it as an optimization problem:

$$\underbrace{arg \min}_H tr(H^T L H) \quad s.t. \quad H^T H = I. \quad (19)$$

Now let's explore Ncut, with objective: $NCut(A_1, A_2, \dots, A_k) = \frac{1}{2} \sum_{i=1}^k \frac{W(A_i, \bar{A}_i)}{vol(A_i)}$, where $vol(A) :=$

$\sum_{i \in A} d_i, d_i := \sum_{j=1}^n w_{ij}$. Similar to RatioCut, we define: $h_{ij} = \begin{cases} 0 & v_i \notin A_j \\ \frac{1}{\sqrt{vol(A_j)}} & v_i \in A_j \end{cases}$. Now

1. Please show that $h_i^T L h_i = \frac{cut(A_i, \bar{A}_i)}{vol(A_i)}$.

$$\begin{aligned} h_i^T L h_i &= \frac{1}{2} \sum_{m=1} \sum_{n=1} w_{mn} (h_{mi} - h_{ni})^2 \\ &= \frac{1}{2} \left[\sum_{m \in A_i, n \notin A_i} \left(w_{mn} \frac{1}{\sqrt{vol(A_i)}} - 0 \right)^2 + \sum_{m \notin A_i, n \in A_i} \left(0 - w_{mn} \frac{1}{\sqrt{vol(A_i)}} \right)^2 + \right. \\ &\quad \left. \sum_{m \in A_i, n \in A_i} \left(w_{mn} \frac{1}{\sqrt{vol(A_i)}} - w_{mn} \frac{1}{\sqrt{vol(A_i)}} \right)^2 + \sum_{m \notin A_i, n \notin A_i} (0 - 0) \right] \\ &= \frac{1}{2} \left[\sum_{m \in A_i, n \notin A_i} \frac{w_{mn}}{vol(A_i)} + \sum_{m \notin A_i, n \in A_i} \frac{w_{mn}}{vol(A_i)} \right] \\ &= \frac{1}{2} \left[\frac{cut(A_i, \bar{A}_i)}{vol(A_i)} + \frac{cut(A_i, \bar{A}_i)}{vol(A_i)} \right] \\ &= \frac{cut(A_i, \bar{A}_i)}{vol(A_i)} \end{aligned} \quad (20)$$

2. Show that $NCut(A_1, A_2, \dots, A_k) = tr(H^T L H)$.

From the definition of Cut,

$$cut(A_1, A_2, \dots, A_k) = \frac{1}{2} \sum_{i=1}^k W(A_i, \bar{A}_i) \quad (21)$$

From the result of Problem 4.1: $h_i^T L h_i = \frac{cut(A_i, \bar{A}_i)}{vol(A_i)}$, we have:

$$\begin{aligned} h_i^T L h_i &= \frac{cut(A_i, \bar{A}_i)}{vol(A_i)} \\ &= \frac{1}{2} \frac{W(A_i, \bar{A}_i)}{vol(A_i)} \end{aligned} \quad (22)$$

So:

$$\begin{aligned} NCut(A_1, A_2, \dots, A_k) &= \frac{1}{2} \sum_{i=1}^k \frac{W(A_i, \bar{A}_i)}{vol(A_i)} \\ &= \sum_{i=1}^k \frac{cut(A_i, \bar{A}_i)}{vol(A_i)} \\ &= \sum_{i=1}^k h_i^T L h_i \\ &= \sum_{i=1}^k (H^T L H)_{ii} \\ &= tr(H^T L H) \end{aligned} \quad (23)$$

3. The constraint now is: $H^T D H = I$.

$$\begin{aligned} h_i^T D h_i &= \sum_j d_l h_{ji}^2 \\ &= \sum_{j \in A_i} d_l \frac{1}{vol(A_i)} \\ &= \frac{1}{vol(A_i)} \sum_{j \in A_i} \sum_{m=1}^N w_{jm} \\ &= 1 \end{aligned} \quad (24)$$

And since $(H^T D H)_{ii} = h_i^T D h_i$, all diagonal elements of $H^T D H$ are 1. Therefore, $H^T D H = \mathbf{I}$.

4. Find the solution to $\underbrace{arg \min}_H tr(H^T L H)$ s.t. $H^T D H = I$.

Define a new matrix, M :

$$M = D^{\frac{1}{2}} H \quad (25)$$

, where D is the degree matrix, and thus $H = D^{-\frac{1}{2}} M$. Then, the minimization problem is equivalent to:

$$\begin{aligned} \min_M Tr(M^T D^{-\frac{1}{2}} L D^{-\frac{1}{2}} M) \\ s.t. M^T M = I \end{aligned} \quad (26)$$

, where $(D^{-\frac{1}{2}})^T = D^{-\frac{1}{2}}$, because the degree matrix D is symmetrical. The problem can be further transformed into:

$$\begin{aligned} \min_M \text{Tr} \left(\frac{M^T D^{-\frac{1}{2}} L D^{-\frac{1}{2}} M}{M^T M} \right) \\ \text{s.t. } M^T M = I \end{aligned} \quad (27)$$

, in which form we can apply the Rayleigh-Ritz theorem, since all matrices in the problem are real matrices ($M^H = M^T$). We know that the minimum value of $\frac{M^T D^{-\frac{1}{2}} L D^{-\frac{1}{2}} M}{M^T M}$ is the smallest eigenvalue of $D^{-\frac{1}{2}} L D^{-\frac{1}{2}}$. Therefore, the solution of M^T is a matrix that has the eigenvectors of L as columns. By substituting $H = D^{-\frac{1}{2}} M$, we can get the solution of H .

Problem 5 [10 pts]

We consider the following optimization problem (\mathbf{Y} is given and generated randomly):

$$\min_{\mathbf{X}} \frac{1}{2} \|\mathbf{X} - \mathbf{Y}\|_F^2 + \|\mathbf{X}\|_* \quad (28)$$

where $\mathbf{Y}, \mathbf{X} \in \mathbb{R}^{100 \times 100}$ and $\|\cdot\|_*$ denotes the nuclear norm (sum of singular values). Now please use gradient descent method to update \mathbf{X} . ($\frac{\partial \|\mathbf{X}\|_*}{\partial \mathbf{X}} = \mathbf{U}\mathbf{V}^T$, where \mathbf{U}, \mathbf{V} is obtained from reduced SVD, namely $[\mathbf{U}, \mathbf{S}, \mathbf{V}] = \text{svd}(\mathbf{X}, 0)$). Plot the objective changes with 1000 iteration.

The objective function is: $L(X, Y) = \frac{1}{2} \|X - Y\|_F^2 + \|X\|_*$, and its derivative w.r.t. X is:

$$\begin{aligned} \frac{\partial L}{\partial X} &= \frac{1}{2}(X - U) + UV^T \\ &= X - Y + UV^T \end{aligned} \quad (29)$$

, where $[U, S, V] = \text{SVD}_{red}(X)$ are derived from the reduced SVD of X .

The change of objective function L w.r.t. number of iterations in gradient descent method is shown in Fig. 1. Codes implemented with MatLab.




Figure 2: Decrease of $L(X, Y) = \frac{1}{2} \|X - Y\|_F^2 + \|X\|_*$ with gradient descent.

Codes (MatLab):

```
Y = 5*rand(100,100);
```

```

X = 5*rand(100,100);
eta = 0.007;
l_L = [];

for i = [1:1000]
if rem(i,100) == 0
i
L(X,Y)
end
l_L = [l_L, L(X,Y)];

X_1 = X - eta * dLdX(X,Y);
X = X_1;
i = i+1;
end

plot(l_L)
xlabel('Iterations')
ylabel('Objective')

%%
function drv = dLdX(X,Y)
[U,S,V] = svd(X,0);
drv = (X - Y) + U * V';
end

function obj = L(X,Y)
obj = 1/2 * norm(X-Y, 'fro')^2 + norm(svd(X,0),1);
end

```

Problem 6 [20 pts]

We turn to Logistic Regression:

$$\min_{\beta} \sum_{i=1}^m \ln(1 + e^{\langle \beta, \hat{x}_i \rangle}) - y_i \langle \beta, \hat{x}_i \rangle, \quad (30)$$

where $\beta = (w; b)$, $\hat{x} = (x; 1)$. Assume $m = 100, x \in \mathbb{R}^{99}$. Please randomly generate x, y and find the optimal β via 1) gradient descent; 2) Newton's method and 3) stochastic gradient descent (SGD) where the batch-size is 1. (need consider choosing appropriate step-size if necessary). Change $m = 1000, x \in \mathbb{R}^{999}$, observe which algorithm will decrease the objective faster in terms of iteration (X -axis denotes number of iteration) and CPU time. [You will receive another 5 bonus points if you implement backtracking line search]

The objective function is $L(\beta, x_i, y_i) = \ln(1 + e^{\langle \beta, \hat{x}_i \rangle}) - y_i \langle \beta, \hat{x}_i \rangle$.

In gradient method, we have the gradient:

$$\frac{\partial L}{\partial \beta} = \frac{e^{\langle \beta, \hat{x}_i \rangle} \hat{x}_i^T}{1 + e^{\langle \beta, \hat{x}_i \rangle}} - y_i \hat{x}_i^T \quad (31)$$

and: $\beta^{i+1} = \beta^i - \lambda \frac{\partial L}{\partial \beta}$. The decrease of L with iterations is shown in the figure below.




Figure 3: Gradient descent ($m=100, n=99$).

In Newton's method, we have:

$$\begin{aligned} \nabla &= \sum_{i=1}^m [p(x_i) - y_i] = X(p - y) \\ H &= XW X^T \end{aligned} \quad (32)$$

, where W is a diagonal matrix with the i -th diagonal element as $p(x_i)[1 - p(x_i)]$. In each update, $w^{i+1} = w^i - H_i^{-1} \cdot \nabla$. The decrease of L with iterations is shown in the figure below.




Figure 4: Newton's method (m=100, n=99).

Note: In Newton's method, the objective function can reach very low values, but will then increase. The performance of Newton's method is not ideal in this case.

In SGD, we have:

$$g = \frac{\partial L}{\partial w} = (p(x_i) - y_i)x_i \quad (33)$$

and for each update $w^{i+1} = w^i - \alpha g$. The decrease of L with iterations is shown in the figure below.




Figure 5: SGD method (m=100, n=99, $\alpha = 0.1$).

From the above results, we can see that the three methods in the order of fastest to slowest convergence is: gradient descent, SGD, Newton's method. The time consumption of the three methods are: gradient descent - 0.01s, Newton's method - 0.29s, SGD - 0.01s. In the case of $m = 1000, n = 999$, the decrease of L with iterations is shown in the figures below.




Figure 6: Gradient descent ($m=1000, n=999$).




Figure 7: Newton's method ($m=1000, n=999$).




Figure 8: SGD method ($m=1000$, $n=999$, $\alpha = 0.1$).

From the above results, we can see that when m, n are increased, the three methods in the order of fastest to slowest convergence is: gradient descent, SGD, Newton's method. The time consumption of the three methods with increased m, n are: gradient descent - 0.04s, Newton's method - 187.81s, SGD - 0.02s.

Codes (MatLab):

```
clear all;

m = 1000;
n = 999;

y = rand(1,m);
w = rand(n,1);
x = [w; rand]
beta = [w; rand];
lambda = 0.1;
l = [];

%%Grd descent
% t = cputime;
% for i = [1:m]
%     x_i = [rand(n,1); 1];
%     y_i = y(i);
%     beta = beta - dLdb(beta,x_i,y_i) * lambda;
%     obj = L(beta,x_i,y_i);
%     l = [l, obj];
% end
```



```

% e = cputime - t;
% disp('GD')
% disp(n)
% disp(e)

%Newton's method

% X = [rand(m,n), ones(m,1)];
% p_vec = zeros(1,m);
%
% t = cputime;
% for i = [1:m]
%     x_i = X(i,:);
%     y_i = y(i);
%
%     px_i = pxi(beta, x_i);
%     obj = L(beta,x_i,y_i);
%     grd = x_i * (px_i - y_i);
%
%     W = zeros(m,m);
%     for i = [1:m]
%         W(i,i)=pxi(beta, X(i,:));
%     end
%     hes = X' * W * X;
%     beta = beta - pinv(hes) * grd';
%
% %     for i = [1:m]
% %         p_vec(i) = pxi(beta, x_i);
% %     end
% %     grd = X' * (p_vec - y)';
% %
% %     W = zeros(m,m);
% %     for i = [1:m]
% %         W(i,i)=pxi(beta, X(i,:));
% %     end
% %     hes = X' * W * X;
% %     beta = beta - inv(hes) * grd;
%
%     l = [l, obj];
% end
% e = cputime - t;
% disp('Newton')
% disp(n)
% disp(e)

```

```

%%SGD

t = cputime;
alpha = 0.1;
for i = [1:m]
y_i = y(i);
x_i = [rand(n,1); 1];
px_i = pxi(beta, x_i);
g = (px_i - y_i) * x_i;
beta = beta - alpha * g;
obj = L(beta,x_i,y_i);
l = [l, obj];
end
e = cputime - t;
disp('SGD')
disp(n)
disp(e)

%% Test

%%
plot(l)
xlabel('Iterations')
ylabel('Objective')
%%
function obj = L(beta,x_i,y_i)
obj = log(1 + exp(dot(beta' , x_i))) - y_i * dot(beta' , x_i);
end

%%
function grd = dLdb(beta,x_i,y_i)
grd = exp(dot(beta' , x_i))/(1 + exp(dot(beta' , x_i))) - y_i * x_i;
end

%%
function likelihood = pxi(beta, x_i)
likelihood = 1/(1+exp(-dot(beta', x_i)));
end

```

Problem 7 [10 pts]

Please design an (either toy or real-world) experiment to demonstrate that PCA can be helpful for denoising.

Take a 480*360 picture of a snowy owl as an example.



Figure 9: Original picture.

Add Gaussian noise with $\mu = 0, \sigma = 0.1$ to all three RGB channels, we have a noisy image.



Figure 10: Noisy picture.

Performing PCA on the noisy picture with the first 100 PCs, we have a reconstructed picture.



Figure 11: Reconstructed picture.

We can see that the noises in the reconstructed picture are smoothed out.

Codes (Mathematica):

```
cwd = NotebookDirectory[];
{m, n} = Import[cwd <> "snowy_owl.jpg", "ImageSize"];
original = Import[cwd <> "snowy_owl.jpg"];
noisy = ImageEffect[original, {"GaussianNoise", 0.1}]
pxl = N[Flatten[ImageData[noisy], 1]];
mean = Mean[pxl];
pxlCtr = # - mean & /@ pxl;
(*RGB channels*)
R = ArrayReshape[pxlCtr[[All, 1]], {m, n}]; G =
ArrayReshape[pxlCtr[[All, 2]], {m, n}]; B =
ArrayReshape[pxlCtr[[All, 3]], {m, n}];

(*Reconstruction single channel*)

recon[channel_, nPC_] := Module[{u, s, v, picrecon},
(*SVD*){u, s, v} = SingularValueDecomposition[channel];
(*Reconstruction using nPC columns*)
picrecon =
u[[All, ;; nPC]].s[[;; nPC, ;; nPC]].Transpose[v[[All, ;; nPC]]];
picrecon]
(*Reconstruct 3 channels*)
```

```

reconRGB[nPC_] :=
ArrayReshape[
Transpose[
Flatten /@ {recon[R, nPC], recon[G, nPC], recon[B, nPC]}], {n, m,
3}]
rgbPlot[plot_] :=
ArrayPlot[ArrayReshape[plot, {n, m, 3}],
ColorFunction -> Function[p, RGBColor[p[[1]], p[[2]], p[[3]]]],
ColorFunctionScaling -> True]

rgbPlot[reconRGB[#]] & /@ {100}

```

Bonus Problem 8 [10 pts]

$$\text{Solve: } \min \|\mathbf{x}\|_0 \quad \text{s.t.} \quad \mathbf{Ax} = \mathbf{y}. \quad (34)$$


We have proved that if $\mathbf{y} = \mathbf{Ax}_o$ with

$$\|\mathbf{x}_o\|_0 \leq \frac{1}{2} \text{krank}(\mathbf{A}). \quad (35)$$

Then \mathbf{x}_o is the unique optimal solution to the ℓ^0 minimization problem

$$\min \|\mathbf{x}\|_0 \quad \text{s.t.} \quad \mathbf{Ax} = \mathbf{y}. \quad (36)$$

However, when \mathbf{A} is of size 5×12 , the following figure illustrates the fraction of success across 100 trials. Apparently $\text{krank}(\mathbf{A}) \leq \text{rank}(\mathbf{A}) \leq 5$, therefore, when sparsity $k = 1, 2$ satisfying Eq. (35)



it has 100% recovery success rate is not surprising. However, the above experiment also shows even $k = 3, 4$ which violates Eq. (35), still it can be recovered at 100%. Please explain this phenomenon.